

Cloud Formation Stack Features :

Step 01: Pre-requisites

- Create default VPC if not present
- Create key pairs
 - Cfn-key-1
 - cfn-key-2
- Gather AMI ID

Step 02: Stack Features

- Create stack
 - Just create an instance without a keypair
- Update stack
 - Add keypair to the above stack
- Create change sets: (same like update but has an additional verification mechanism)
 - Change the key-1 to key-2
 - Change the instance type
- Rollback
 - Change the keypair to key-3 which is non existing one.

Resources : (Required)

- Resources are the key components of a CF stack.

Syntax:

Resources:

Logical ID:

Type: Resource Type

Properties:

Set of properties

Example:

Resources:

MyEC2Instance:

Type: "AWS::EC2::Instance"

Properties:

ImageID: "ami-1234456ff123"

!Ref : Intrinsic Function

- Ref returns the value of the specified parameter or resource, so basically it can be used in parameter or resource section of the CF.

- **Resource Case:** When we specify a resource logical name, it returns a value that we can typically use to refer to that resource.
- **Parameter Case :** When we specify a parameter logical name, it returns the value of that parameter.

Syntax :

- Long Form:
 - Ref: logicalName
- Short Form:
 - !Ref logicalName

Example:

```
MyEIP:
  Type: "AWS::EC2::EIP"
  Properties:
    InstanceId: !Ref MyEC2Instance
```

Parameters:

- It enables us to input custom values to our template each time when we create or update a stack.
- Maximum 60 parameters can have a single CFN template.
- Must be given a logical name (logical ID) which must be alphanumeric and unique among all logical names within the template.
- Each parameter must be assigned a parameter type that is supported by AWS CF.
- Each parameter must be assigned a value at runtime for AWS CF to successfully provision the stack. We can optionally specify a default value to unless another value is provided.
- Must be declared and referenced within the same template.
- We can reference parameters from the resources and outputs sections of the template.

Syntax:

```
Parameters:
  ParameterLogicalId:
    Type: DataType
    ParameterProperty: Value
```

Examples:

```
Parameters:
  InstanceTypeParameter:
    Type: String
    Default: t2.micro
    AllowedValues:
      - t2.micro
      - m1.small
      - m1.large
```

Description: Please enter the required Instance type.

Parameter Properties:

- AllowedPattern
- AllowedValues
- ConstraintDescription
- Default
- Description
- MaxLength
- MaxValue
- MinLength
- MinValue
- NoEcho

Parameter Types:

- String
- Number
- List
- CommaDelimitedList
- AWS Specific
 - AWS::EC2::Instance::Id
 - AWS::EC2::VPC::Id
 - List<AWS::EC2::Subnet::Id>
- SSM Parameter Type
 - AWS::SSM::Parameter::Name
 - AWS::SSM::Parameter::Value<String>
 - AWS::SSM::Parameter::Value<List<String>>

Mappings:

- Matches a key to a corresponding set of named values.
- For an example, if we wanted to set values based on a region, we can create a mapping that uses region name as a key and contains the values we want to specify for each region.
- We can use Fn::FindInMap intrinsic function to retrieve values in map

Syntax:

Mappings:

Mapping01:

Key01:

Name: Value01

Key02:

Name: value02

Key03:

Name: Value03

Examples:

Mappings:

RegionMap:

us-east-1:

"HVM64": "ami-11111ddd111"

eu-west-1:

"HVM64": "ami-334323ffd222"

ap-southeast-1:

"HVM64": "ami-9767ghhh1122"

Intrinsic Function: FindInMap :

- FindInMap returns the value corresponding to keys in a **two-level map** that is declared in mappings section.
- **Parameters:**
 - Map Name
 - Top Level Key
 - Second Level Key
 - Return Value

Example:

Mappings:

RegionMap:

us-east-1:

"HVM64": "ami-11111ddd111"

"HVMG2": "ami-5735ggh445"

eu-west-1:

"HVM64": "ami-334323ffd222"

"HVMG2": "ami-557556474t5"

ap-southeast-1:

"HVM64": "ami-9767ghhh1122"

"HVMG2": "ami-96468vdfhsh222"

Resources:

DevEC2Instance:

Type: "AWS::EC2::Instance"

Properties:

ImageId: !FindInMap

- RegionMap

- !Ref 'AWS::Region'

- HVM64

Pseudo Parameters:

- These parameters are predefined by AWS CF.
- Don't need to declare them in our template.
- We can use them the same way as we use parameters as an argument for **Ref** function.
- Usage:
 - Outputs:
 - DevStacksRegion:
 - Value: !Ref "AWS::Region"

Parameters:

- AWS::AccountId
- AWS::NotificationARNs
- AWS::NoValue
- AWS::Partition
- AWS::Region
- AWS::StackId
- AWS::StackName
- AWS::URLSuffix

Conditions:

- It contains statements that define the circumstances under which entities are created or configured.
- **Exmpl-1:** we can create a condition and then associate it with a resource or output so that AWS CF only creates the resource or output if the condition is true.
- **Exmpl-2:** we can associate the condition with a property so that AWS CF only sets the property to a specific value if the condition is true, if the condition is false, AWS CF sets the property to a different value that we specify.
- We will use conditions, when we want to reuse the template in different contexts like dev and prod environments.
- Within each condition we can reference the other condition.
- We can associate these conditions in three places.
 - Resources
 - Resource Properties
 - Outputs
- A stack creation or stack update, AWS CF evaluates all conditions in the template. During stack update, resources that are now associated with a false condition are deleted.

Syntax:

Conditions:
Logical ID:
Intrinsic function

Example:

Conditions:

```
CreateEIPForProd:
  Fn::Equals:
    - !Ref EnvironmentName
    - Prod
```

Conditions - Intrinsic Functions :

- Below are the list of conditions which can be used in AWS CF.
 - Fn::And
 - Fn::Equals
 - Fn::If
 - Fn::Not
 - Fn::Or

Outputs: (Optional)

- It declare output values that we can use as below
 - Import into another stack (to create cross-stack references)
 - When using Nested Stacks we can see how outputs of a nested stack are used in Root Stack.
 - We can view outputs on the CF Console.
 - We can declare maximum of 60 outputs in a CF template

Syntax:

Outputs:

Logical ID:

Description: Information about the value

Value: Value to return

Export:

Name: Value to export

Example:

Outputs:

InstanceId:

Description: Instance ID

Value: !Ref MyVMInstance

Export:

Name: !Sub "\${AWS::StackName}-InstanceId"

MyInstAvailabilityZone:

Description: Instance Availability Zone

Value: !GetAtt MyVMInstance.AvailabilityZone

Export:

Name: !Sub "\${AWS::StackName}-InstanceAz"

Export (Optional) :

- Contains resource output used for cross stack reference.
- For each AWS account Export name must be unique within the region. As it should be unique we can use the export name as **"AWS::StackName"-ExportName**
- We can't create cross stack references across region.
- We can use the intrinsic function Fn::ImportValue to import value that have been exported within the same region.
- For outputs the value of the name property of an export can't use Ref or GetAtt functions that depend on a resource.
- We can't delete a stack if another stack references one of its output.
- We can't modify or remove an output value that is referenced by another stack.
- We can use outputs in combination with Conditions.

Metadata:

- Metadata provides details about the cfn template.

Syntax:

Instances:

Description: "Info about the Instances"

Databases:

Description: "Info about the Databases"

Example:

AWSTemplateFormatVersion: 2010-09-09

Metadata:

Instances:

Description: My VM Instance

- We have three types of metadata keys which are listed below
 - AWS::CloudFormation::Designer
 - Auto generated during resources drag and drop to canvas
 - Designer, visually depicts how our resources are laid out.
 - Automatically add information when we use it to create, view and update templates. Its a system generated metadata.
 - Defines the info about the resources such as their Size and Relative position in template metadata. All layout info is stored in designer metadata.
 - We drag and drop the resources.
 - When we create template,It enforces some basic relationships between resources to help us create valid template.
 - Example: we can not directly add EC2 Instances in VPC, we must add a subnet in a VPC.
 - We can validate templates directly in designer.

- We can validate our manually written template using Designer validate template.
 - Conversion of JSON to YAML or vice versa is possible in designer.
- AWS::CloudFormation::Interface
 - Used for parameter grouping
 - When we create or update stacks in the console, the console lists input parameters in alphabetical order by their logical IDs.
 - By Using this Key, we can define our own parameter grouping and ordering so that users can efficiently specify parameter values.
 - We can also define labels for parameters.
 - A label is friendly name or description that the console displays instead of a parameter's logical ID which helps users understand the values to specify for each parameter.

Syntax:

Metadata:

AWS::CloudFormation::Interface

ParameterGroups:

- ParameterGroup

ParameterLabels:

- ParameterLabel

Example:

Metadata:

AWS::CloudFormation::Interface

ParameterGroups:

- Label:

default: "EC2 Instance configuration"

Parameters:

- InstanceType

- KeyName

- Label:

Default: "Environment Configuration"

Parameters:

- EnvironmentName

ParameterLabels:

EnvironmentName:

Defaults: "which env we are planning to"

- AWS::CloudFormation::Init
 - Used for application installation and configurations on AWS Compute EC2 instances

EC2UserData:

- We can use User Data in CF template for EC2.

- We need to use intrinsic function Fn::Base64 with UserData in CFN template. This function returns the Base64 representation of input string. It passes encoded data to EC2 Instance.
- YAML Pipe (|) : Any intended text that follows should be interpreted as a multiline scalar value which means value should be interpreted literally in such a way that preserves newlines.
- UserData Cons:
 - By default UserData scripts and cloud-init directives run only during the boot cycle when we first launch an instance.
 - We can update our configuration to ensure that our userdata scripts and cloud init directives run every time we restart our instance. (Reboot of server required)

Sample:

UserData:

Fn::Base64:

```
#!/bin/bash
Sudo yum update
Sudo yum -y erase java-1.7.0-openjdk.x86_64
Sudo yum -y install java-1.8.0-openjdk.x86_64
Sudo yum -y install java-1.8.0-openjdk-devel
Sudo yum -y install tomcat8
Service tomcat8 start
Mkdir /usr/share/tomcat8/webapps/ROOT
Touch /usr/share/tomcat8/webapps/ROOT/index.html
Echo "Hello World" > /usr/share/tomcat8/webapp/ROOT/Index.html
```